Agent 365

From Chaos to Control: Managing AI Agents at Scale with Agent 365

Kranthi Manchikanti
AI Solutions Architect

Microsoft

Background

Copilots vs Agents

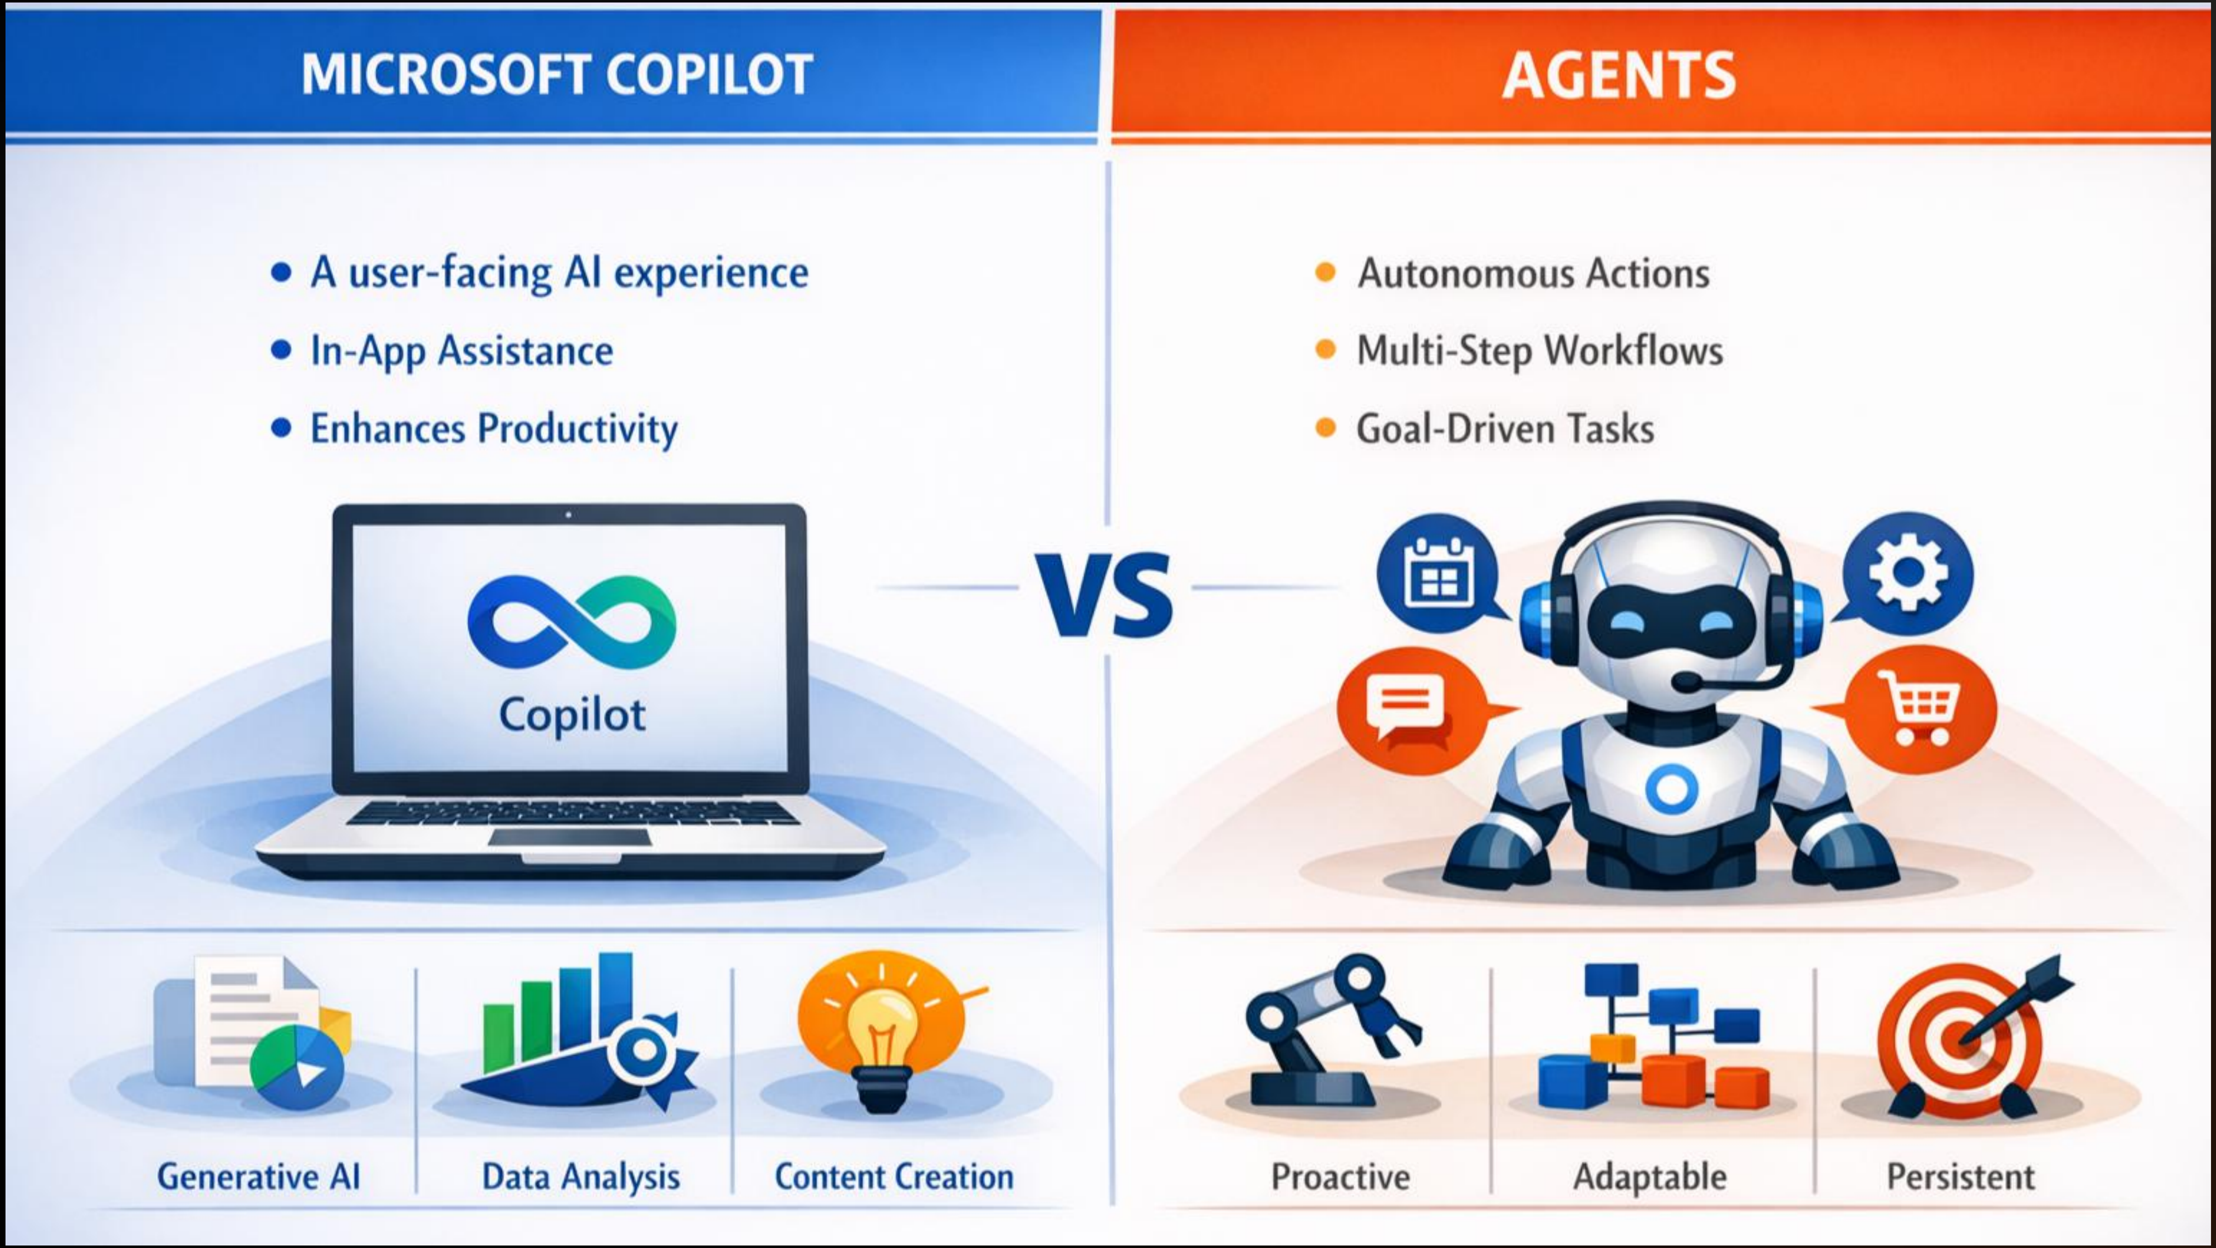

Microsoft 365 Copilot architecture

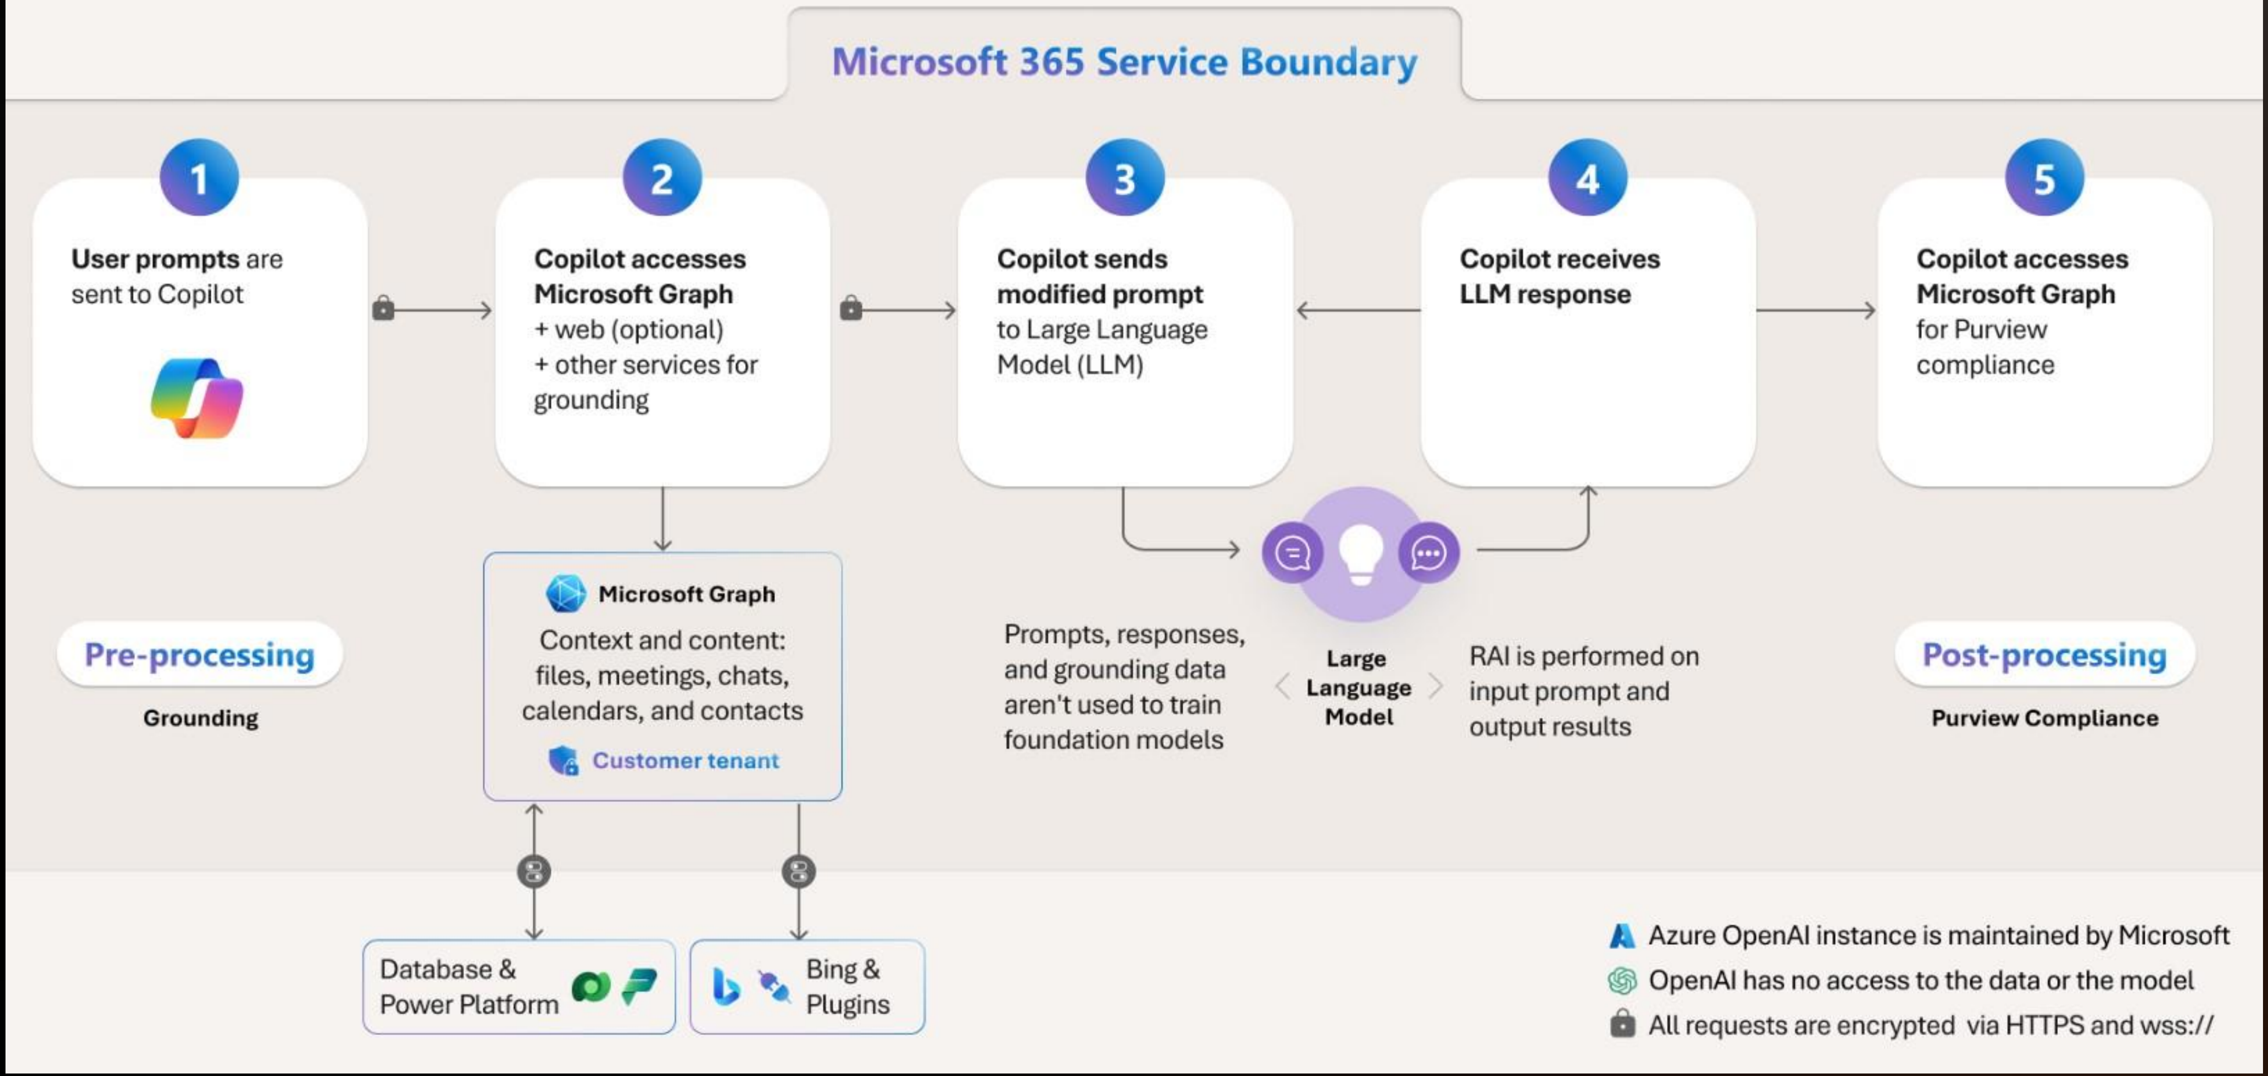

What is Agentic AI?

Agentic AI refers to autonomous AI systems that plan, reason, and act to complete tasks with minimal human input.

Agents use AI to automate and execute business processes, working alongside or on behalf of a person, team, or organization.

Deploy to...

- Microsoft Copilot
- Your applications
- Your websites

Key Characteristics

- **Goal-oriented:** Take goal-directed actions with minimal human oversight
- **Context Understanding:** Comprehend and follow natural language and other modalities.
- **Multi-Step Reasoning:** Provide contextual decision-making and make judgment calls and tradeoffs.
- **Adaptable Planning:** Dynamically adjust plans based on changing conditions to complete processes.
- **Action Enabled:** Empowered to act via access to web services to deliver specific skills.

Agent internals

"System that uses a LLM to decide the control flow of an application."

Autonomy Levels:

- **No Autonomy:** Traditional RAG
- **Simple:** Paths routing
- **Fully:** Multi-step reasoning & acting

Architectures:

- Single agent
- Multi agent

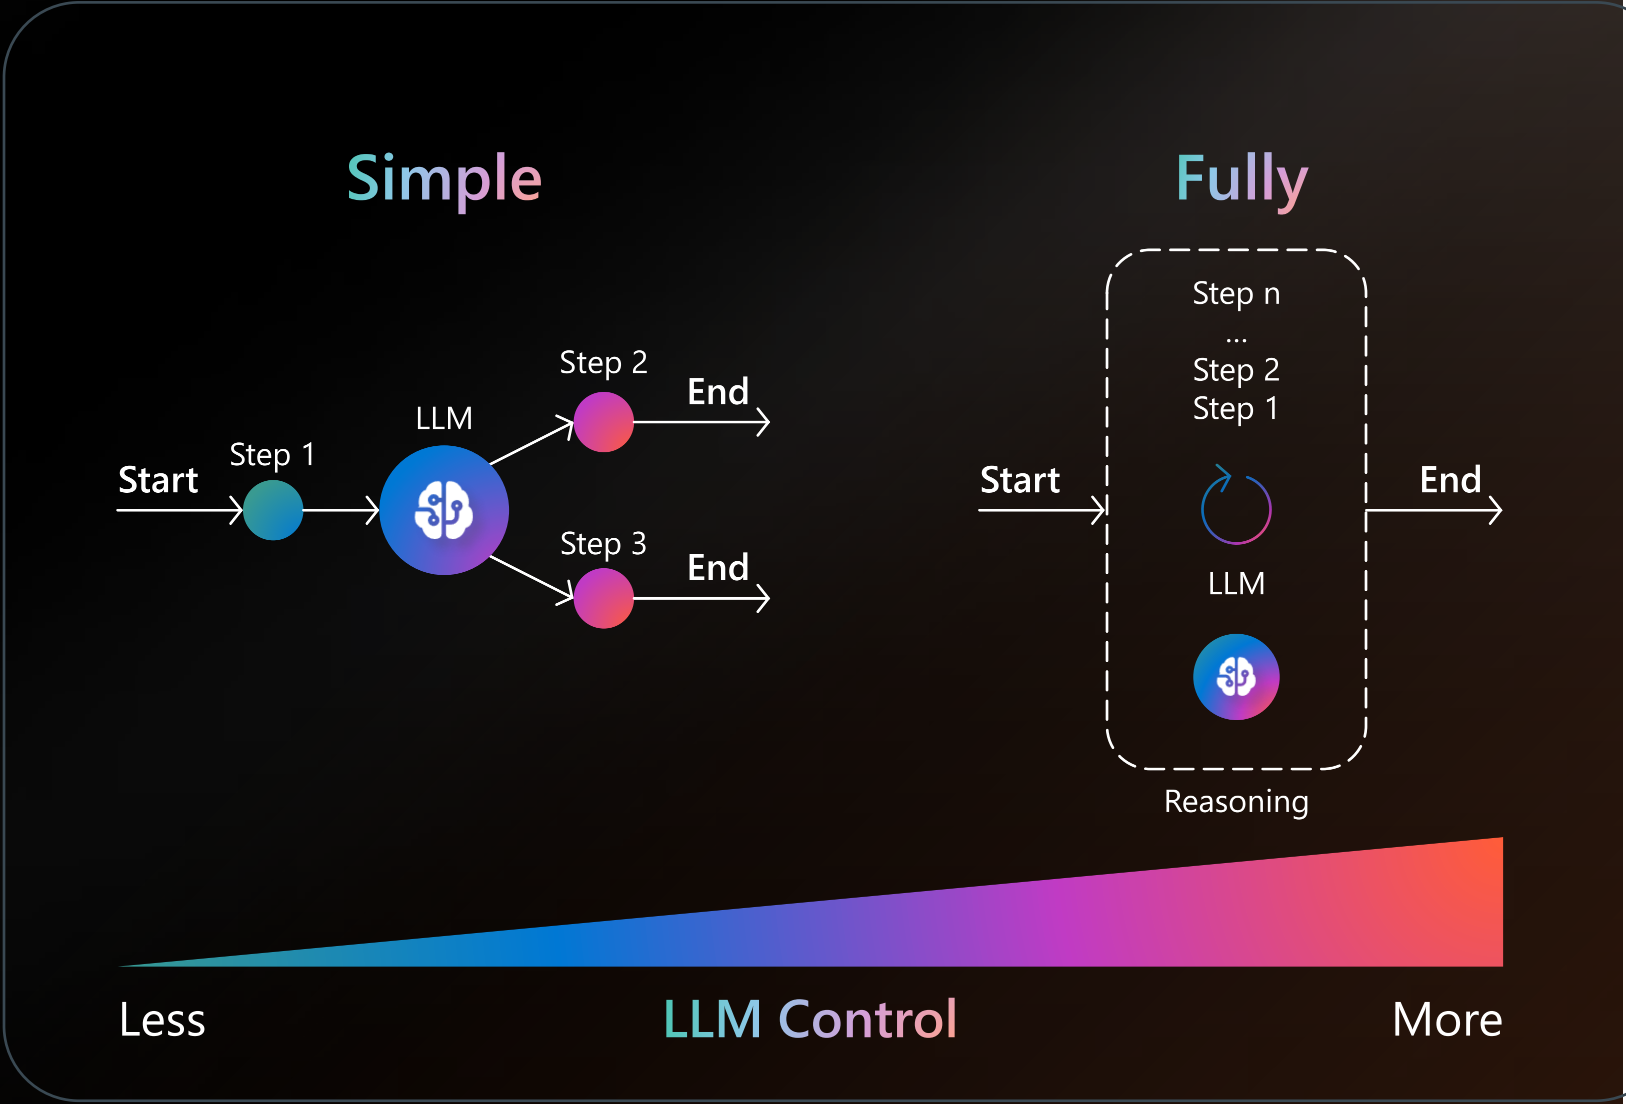

M365

People + Copilot + Agents

Agents are joining the workforce

1.3 Billon : Projected Number of agents by 2028

Source: IDC Info Snapshot, 1.3 Billion AI Agents by 2028, doc #US53361825, May 2025

Every organization is transforming into a frontier firm

Pattern 1

Human with assistant

Pattern 2

Human-agent teams

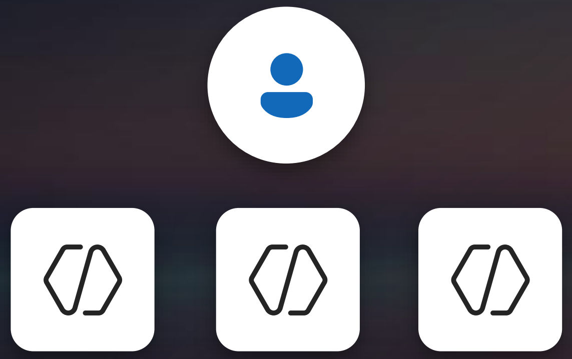

Pattern 3

Human-led, agent operated

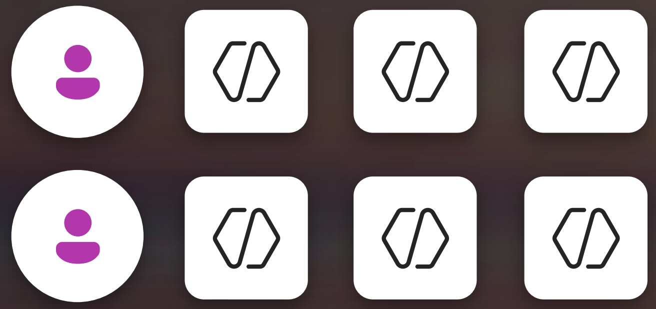

Is your organization ready?

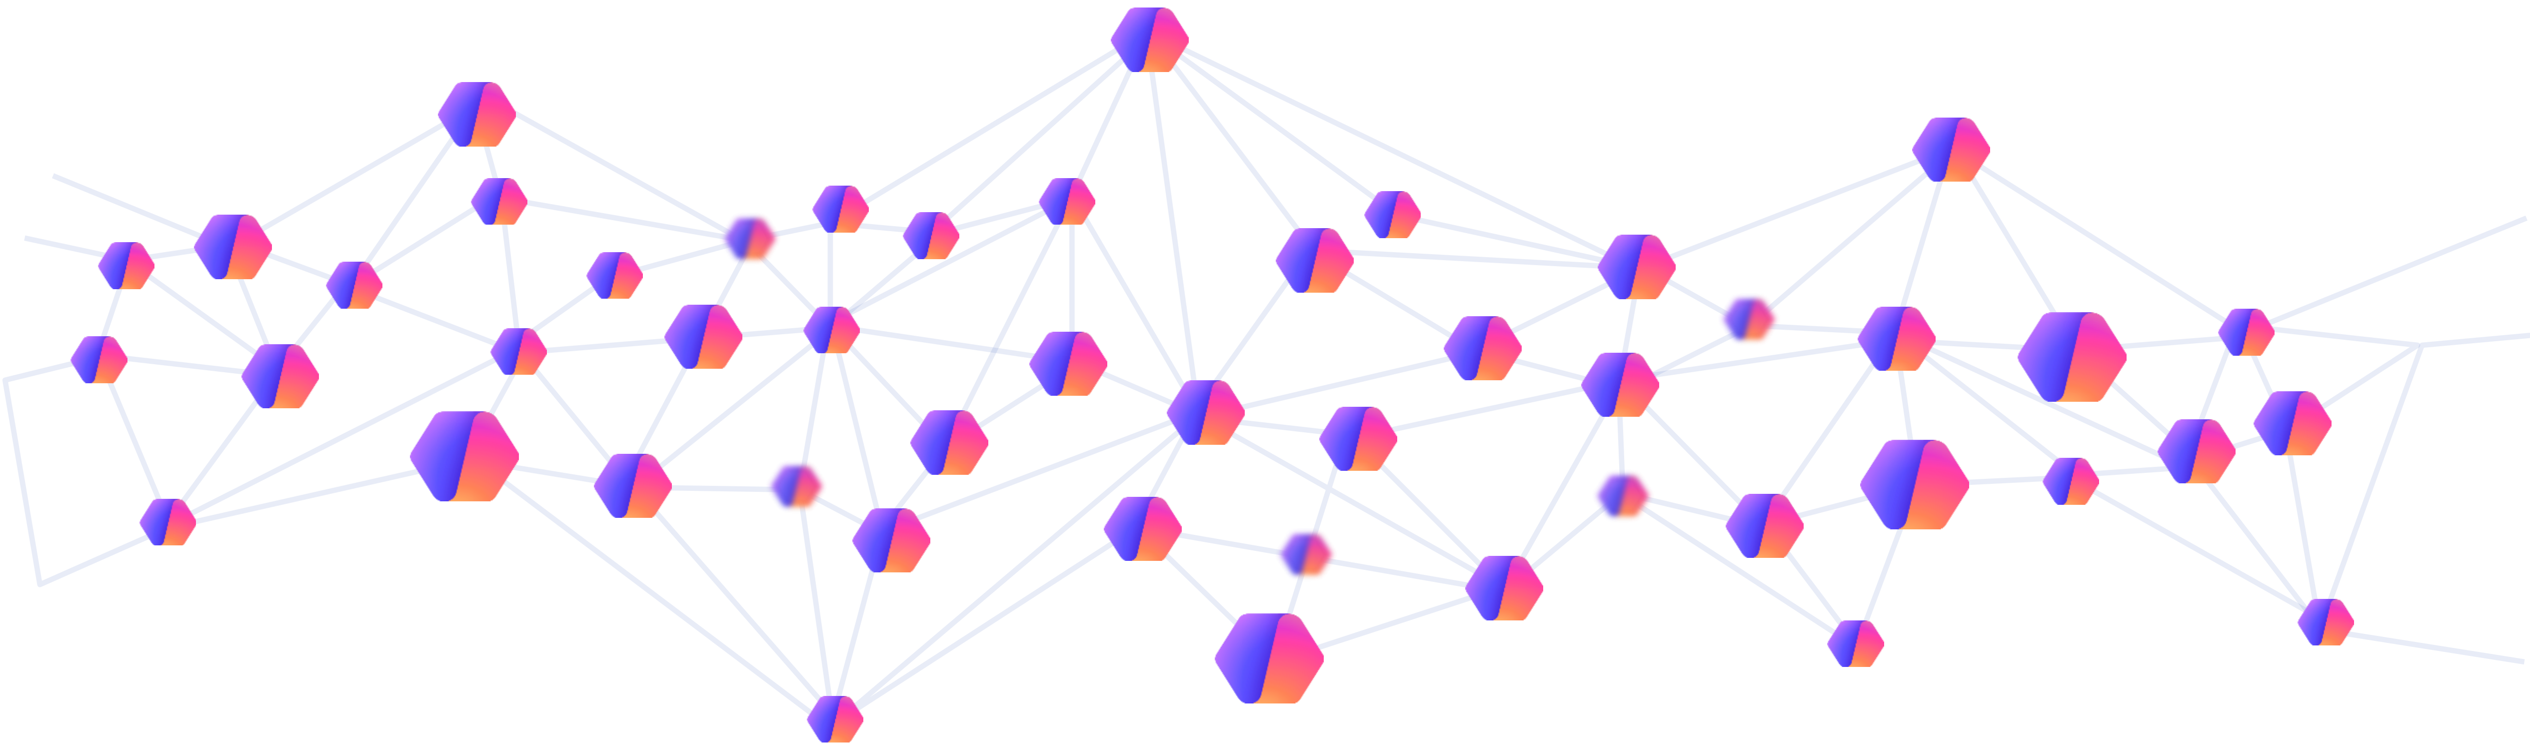

Can IT discover and manage agent proliferation for impact?

Are agents behaving correctly & appropriately within the enterprise?

Who / what are agents sharing the sensitive information with?

Are the agents well governed and audited – what are my costs?

AI Adoption Needs for Enterprises

Inventory and observability of all agents

Consistency across all agent runtimes

Trust, governance and security

Clarity on agent ROI and outcomes

Readiness for new workflows and roles

Introducing

Microsoft Agent 365

The control plane for agents

Why is Agent 365 Needed?

- **Organizations are scaling AI rapidly**
  - Most companies now run dozens or even hundreds of AI automations, copilots, and agents across teams. Agent 365 brings standardization and avoids scattered, unmanaged AI.
- **Ensures security and compliance at scale**
  - Agent 365 provides:
    - Centralized permissions
    - Built-in compliance
    - Unified auditing
- **Reduces risk of unmanaged or "shadow" agents**
  - Agent 365 detects and governs:
    - Agents built internally
    - Agents purchased from vendors
    - Agents that appear organically in workflows

Enables cross-application and multi-agent workflows

- Agents can collaborate, exchange information, and execute tasks across:
  - Teams
  - SharePoint
  - Calendars

Simplifies lifecycle management

- IT teams can (All from one place):
  - Register and certify agents
  - Assign owners
  - Track usage and performance

Built for enterprise-grade identity and governance

- Microsoft emphasizes that Agent 365 is for agents just as M365 is for human users—meaning organizations can secure and monetize AI workloads separately and safely.

What is Agent 365

"Microsoft Agent 365 is Microsoft's enterprise control plane for managing AI agents across Microsoft 365 and third-party systems. It provides a unified layer to build, deploy, govern, secure, monitor, and scale AI agents in a consistent and compliant manner."

Key Capabilities

1. Centralized Registry & Fleet View
2. Identity & Access Control
3. Enterprise-Grade Governance
4. Observability & Telemetry
5. Developer Enablement
6. Integration with Copilot & Agent Mode

An IT Control Plane for all of your agents

Microsoft Agent 365

Microsoft Agent 365 – Foundational Elements

Enterprise-Grade AI Agent Control

Agent 365 treats AI agents as enterprise assets with identities, policies, and lifecycle management.

Integration with Microsoft Ecosystem

Seamlessly integrates with Microsoft 365, Entra, Defender, and Purview for unified governance.

Centralized Management and Visibility

Provides a single pane of glass for monitoring agent presence, access, behavior, and value.

Governance and Security Compliance

Ensures governance, traceability, and protection aligned with Zero Trust and compliance principles.

Microsoft Agent 365 - Observability

Agent Registry

Central inventory of all AI agents (Microsoft, third-party, custom) with ownership and status

Agent Map

Visual view of agent relationships, dependencies, and interactions with data, tools, and other agents

Agent Analytics

Metrics on usage, performance, quality, and business impact

Role-Specific Oversight

Tailored views for IT, Security, Risk, and Business leaders

Microsoft Agent 365 - Governance

Agent Onboarding

Controlled registration, approval, and ownership assignment

Integration Management

Govern how agents connect to apps, APIs, data sources, and other agents

Lifecycle Management

Policies for updates, inactivity, ownership changes, and retirement

Audit & Logging

Immutable activity logs for investigations and regulatory reviews

Data Compliance

Enforcement of organizational and regulatory data policies

Microsoft Agent 365 - Security

Access Control

Identity-based access using Microsoft Entra Agent ID and Conditional Access

Data Security

Protection against oversharing, data leakage, and unsafe content usage

Threat Protection

Detection and response for malicious or compromised agent behavior via Microsoft Defender

Microsoft Agent 365 SDK and CLI

> ⓘ Important
>
> You need to be part of the Frontier preview program to get **early access** to Microsoft Agent 365. Frontier connects you directly with Microsoft's latest AI innovations. Frontier previews are subject to the existing preview terms of your customer agreements. As these features are still in development, their availability and capabilities may change over time.

Agent 365 SDK

Use the **Agent 365 SDK** to extend agents built using any agent SDK or platform, with enterprise-grade identity, observability, notifications, security, and governed access to Microsoft 365 data.

Agents have unique identities. People invoke them using common gestures (such as @mentions) in apps that enterprise users typically operate in (such as Teams, Word, Outlook, and more). They demonstrate observable behaviors that build trust, take auditable actions, and do so via secure access to tools and data.

With the Agent 365 SDK, agents can:

- Use Entra-backed Agent Identity with their own user resources like mailbox for secure authentication and controlled access to tools and data.
- Receive and respond to notifications from Teams, Outlook, Word comments, and emails—just like a human participant in Microsoft 365 apps.
- Gain full observability via Open Telemetry, enabling audited, traceable agent interactions, inference events, and tool usage.
- Invoke governed Model Context Protocol (MCP) servers to access Microsoft 365 workloads (for example, Mail, Calendar, SharePoint, Teams) under admin control.
- Function within an IT-approved blueprint system, ensuring each agent instance inherits compliance, governance, and security policies.

Learn more about the Agent 365 SDK.

All your agents can be enabled for Agent 365

Agent 365

- Agent Registry and Access
- Agent Visualization and Security
- Agent Interoperability

+

Your Agent

- Copilot Studio
- Microsoft Foundry
- Microsoft Agent Framework
- SDK
  Any AI Stack + Agent 365 SDK

Demo

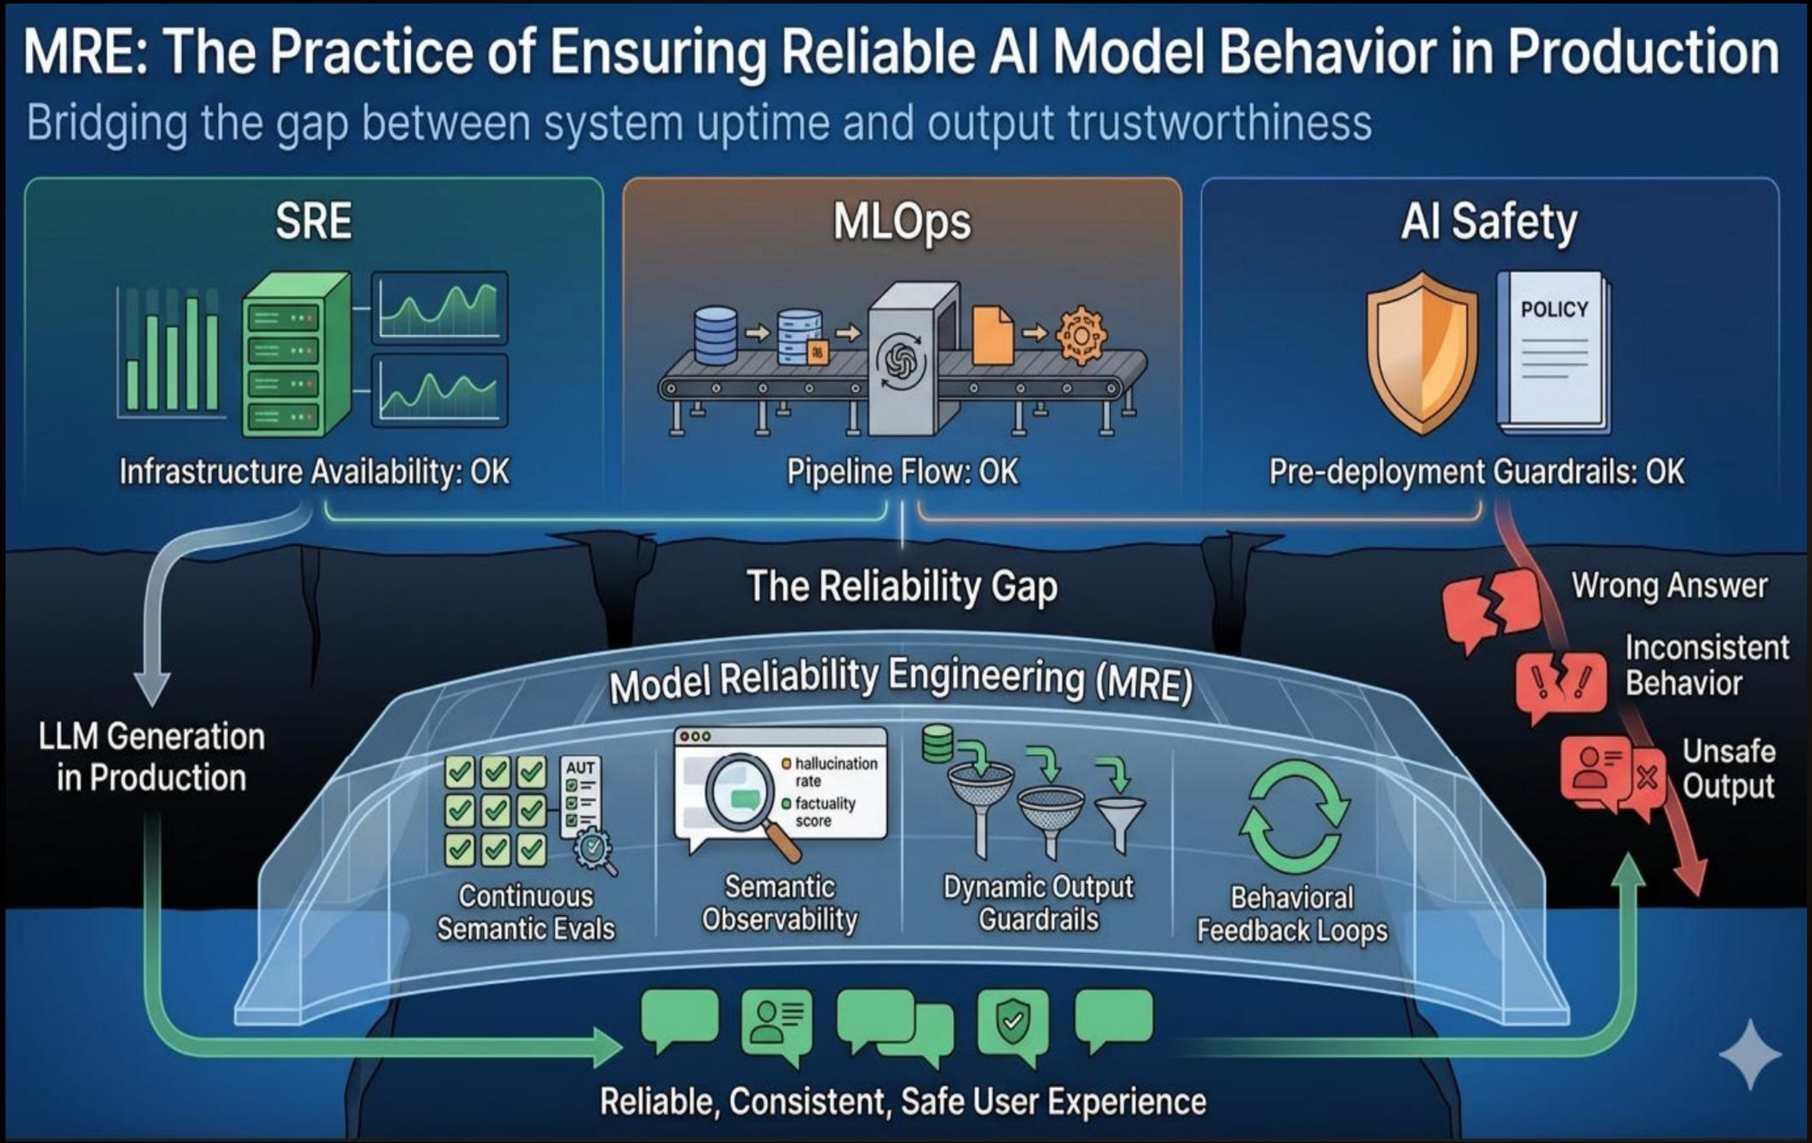

(1) Model Reliability Engineering: Who Owns It When the AI Is Confidently Wrong?

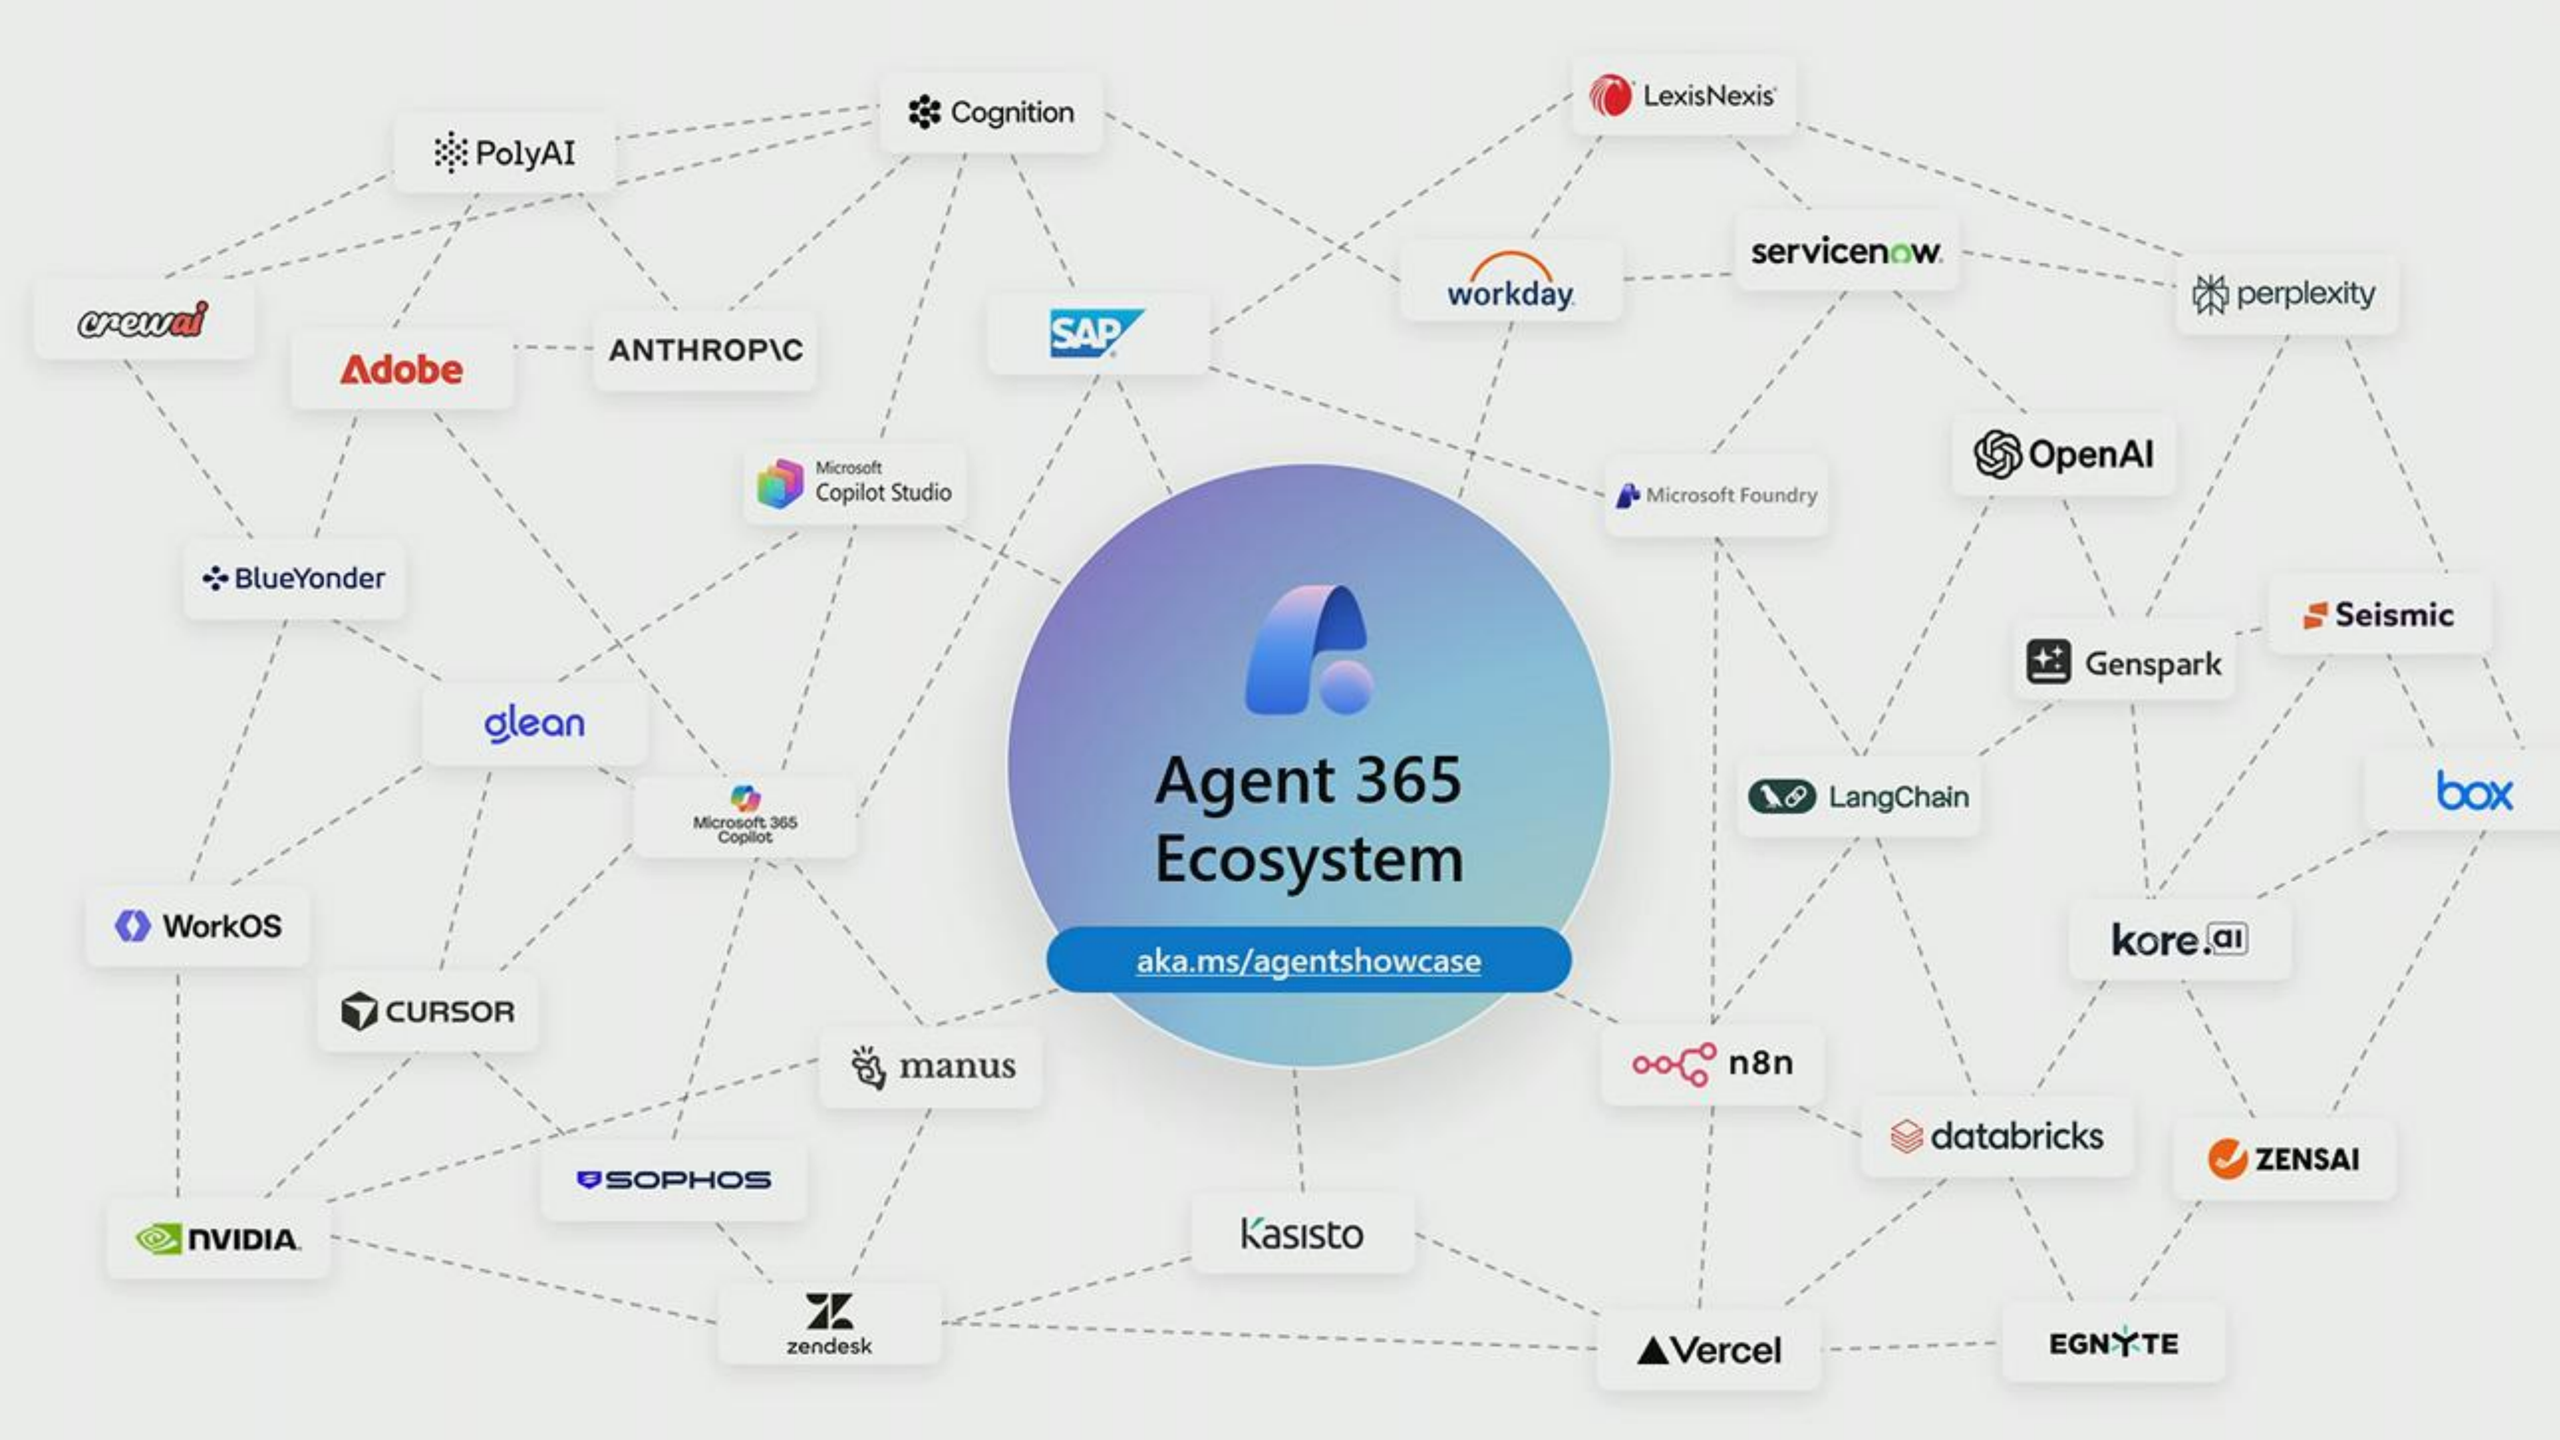
Agent 365
Ecosystem
aka.ms/agentshowcase
Cognition
PolyAI
crewai
Adobe
ANTHROP\C
SAP
LexisNexis
workday
servicenow
perplexity
Microsoft Copilot Studio
Microsoft Foundry
OpenAI
BlueYonder
Seismic
Genspark
glean
Microsoft 365 Copilot
LangChain
box
WorkOS
kore.ai
CURSOR
manus
n8n
databricks
ZENSAI
SOPHOS
NVIDIA
kasisto
zendesk
Vercel
EGNYTE

Thank you!